

Machine Learning Platform Engineer
Helio Labs — London

ABOUT US

Helio Labs trains and serves large models for scientific research groups.

WHAT YOU'LL DO

- Build and operate distributed training infrastructure on GPU clusters.
- Optimise model serving latency and throughput for inference workloads.
- Work directly with research scientists to productionise experimental models.
- Own our feature store and the pipelines that populate it.

REQUIREMENTS

- Strong Python, with production PyTorch experience.
- Experience with distributed training: data and model parallelism.
- CUDA familiarity and GPU performance profiling.
- Experience with an ML orchestration framework such as Kubeflow or Ray.
- Understanding of vector databases and embedding retrieval at scale.
- MLOps practice: experiment tracking, model registries, reproducibility.

NICE TO HAVE

- Publications at a major ML venue.
- Experience with quantisation or model distillation.
- Familiarity with Triton Inference Server.

BENEFITS

- Competitive salary and equity.
- Compute budget for personal research.
- Conference attendance, including one international conference a year.

EQUAL OPPORTUNITY

Helio Labs is proud to be an equal opportunity workplace and is an affirmative action employer.